Hands on Introduction to
IBM Data Science Experience

Power of data. Simplicity of design. Speed of innovation.

Asad Mahmood
Jean Bright
Charles Morrison

Agenda

Overview of DSX, Watson Data Platform and Use Case (Wednesday Feb. 28, 2018)

3:00 PM - 4:00 PM	**Introduction to IBM Data Science Experience and Watson Data Platform**
4:00 PM - 5:00 PM	**Overview of Use Case and Solution Approach**

Hands-On DSX Labs (Thursday March 1, 2018)

9:00 AM - 10:30 AM	**Lab 1 - Setting Up Your DSX Environment and Exploratory Data Analysis**
10:30 AM - 12:00 PM	**Lab 2 - Data Visualization with R Studio and Shiny**
12:00 PM - 12:30 PM	**Lunch**
12:30 PM - 1:30 PM	**Lab 3 - Building a Predictive Model with Watson Machine Learning**
1:30 PM - 2:00 PM	**Next steps regarding POC and project timeline**

Participant Background

Open Source

- R/Python/Scala
- Jupyter Notebook
- Spark
- Hadoop

IBM

- Bluemix
- Data Science Experience
- Watson Machine Learning

What is Data Science?

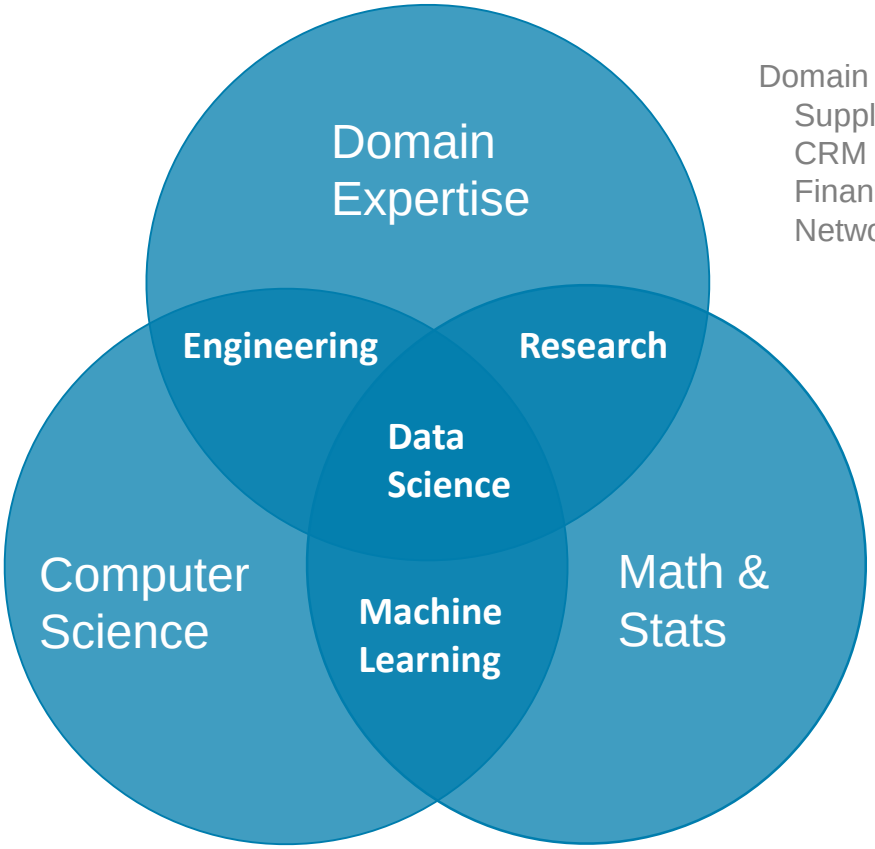

Data Science Projects Require Multiple Skills

Data Science Impact Across Industries and Use Cases

$10s of Billions in each industry and use case

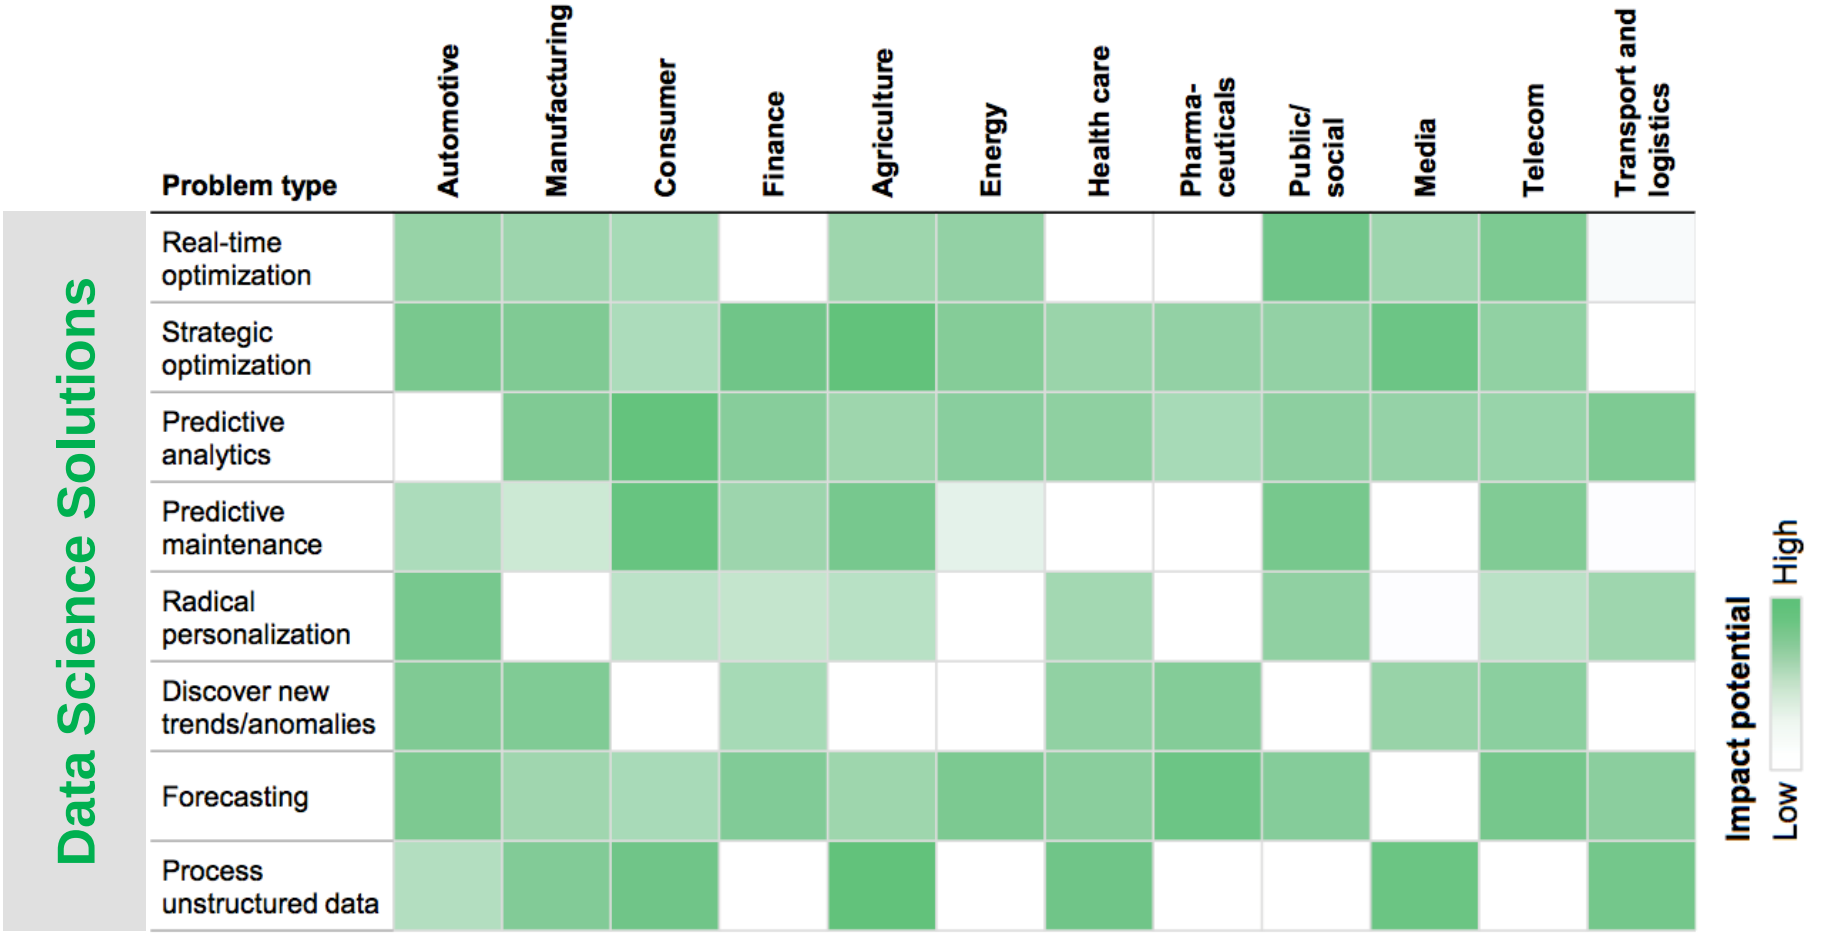

SOURCE: McKinsey Global Institute analysis

Challenges in delivering value with Data Science

Data

- Data resides in silos and difficult to access
- Detailed data was never stored
- Unstructured and external data wasn't considered

Governance

- Self-service isn't a reality, if the data isn't secure
- Understanding lineage and getting to a system of truth

Skills

- Data Science skills are in low supply and high demand
- Nurturing new data professionals is challenging

Infrastructure

- Need an environment that enables collaboration and deployment to production
- Discrete tools present barriers to progress

Watson Data Platform

IBM Watson Data Platform

Mission: Make Data Simple and Accessible to All

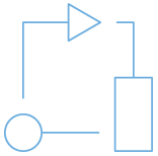

Platform.

Method.

Ecosystem.

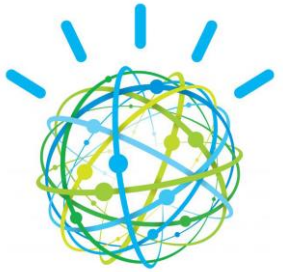

http://ibm.co/makedatasimple

IBM Watson Data Platform

Experience New Ways To Put Data To Work

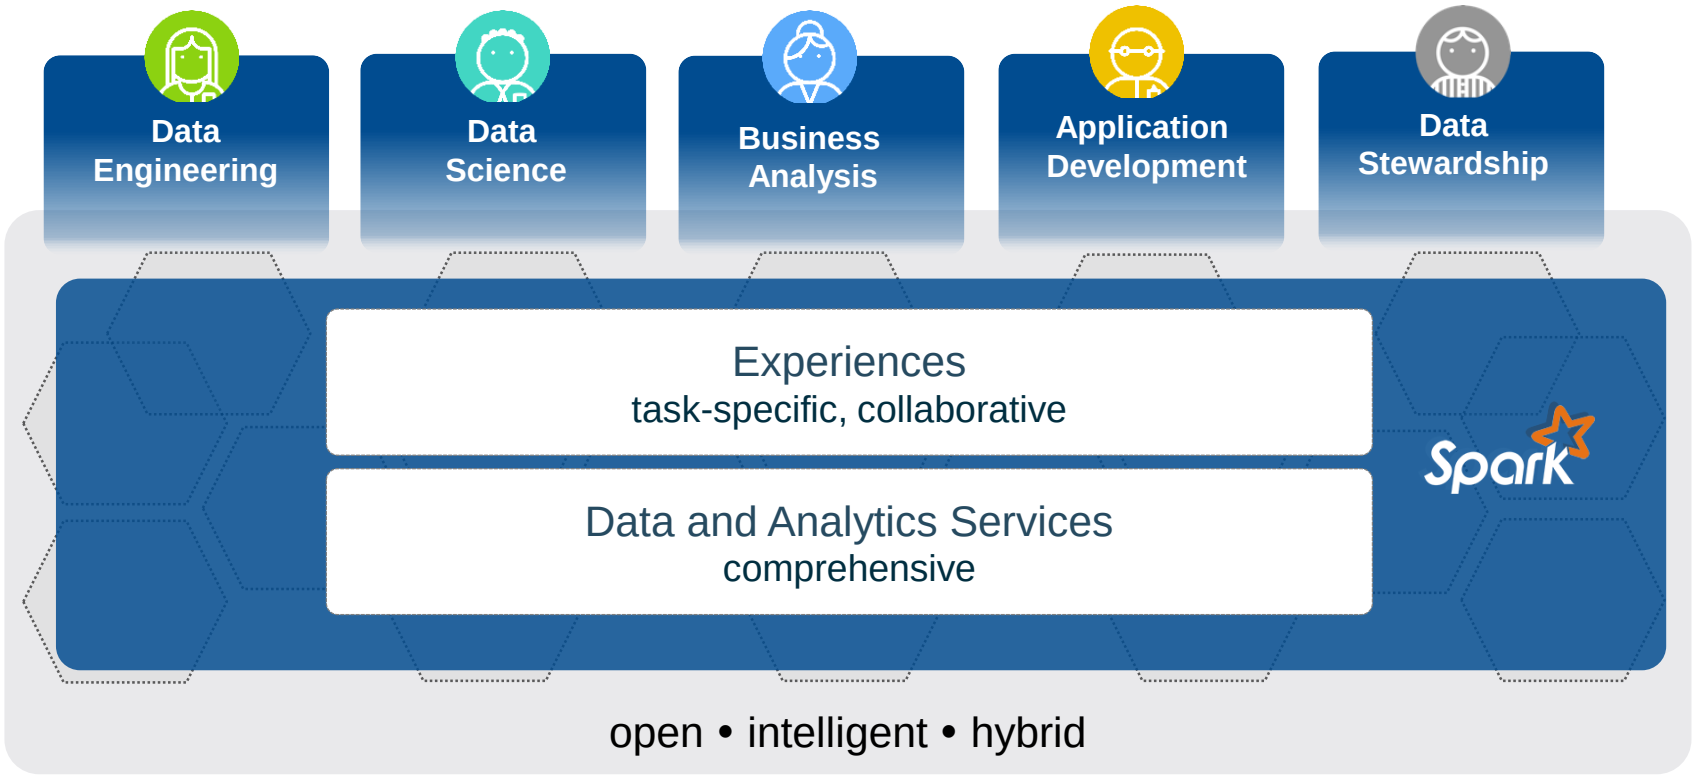

IBM Watson Data Platform

Connects Users to Data and Analytics

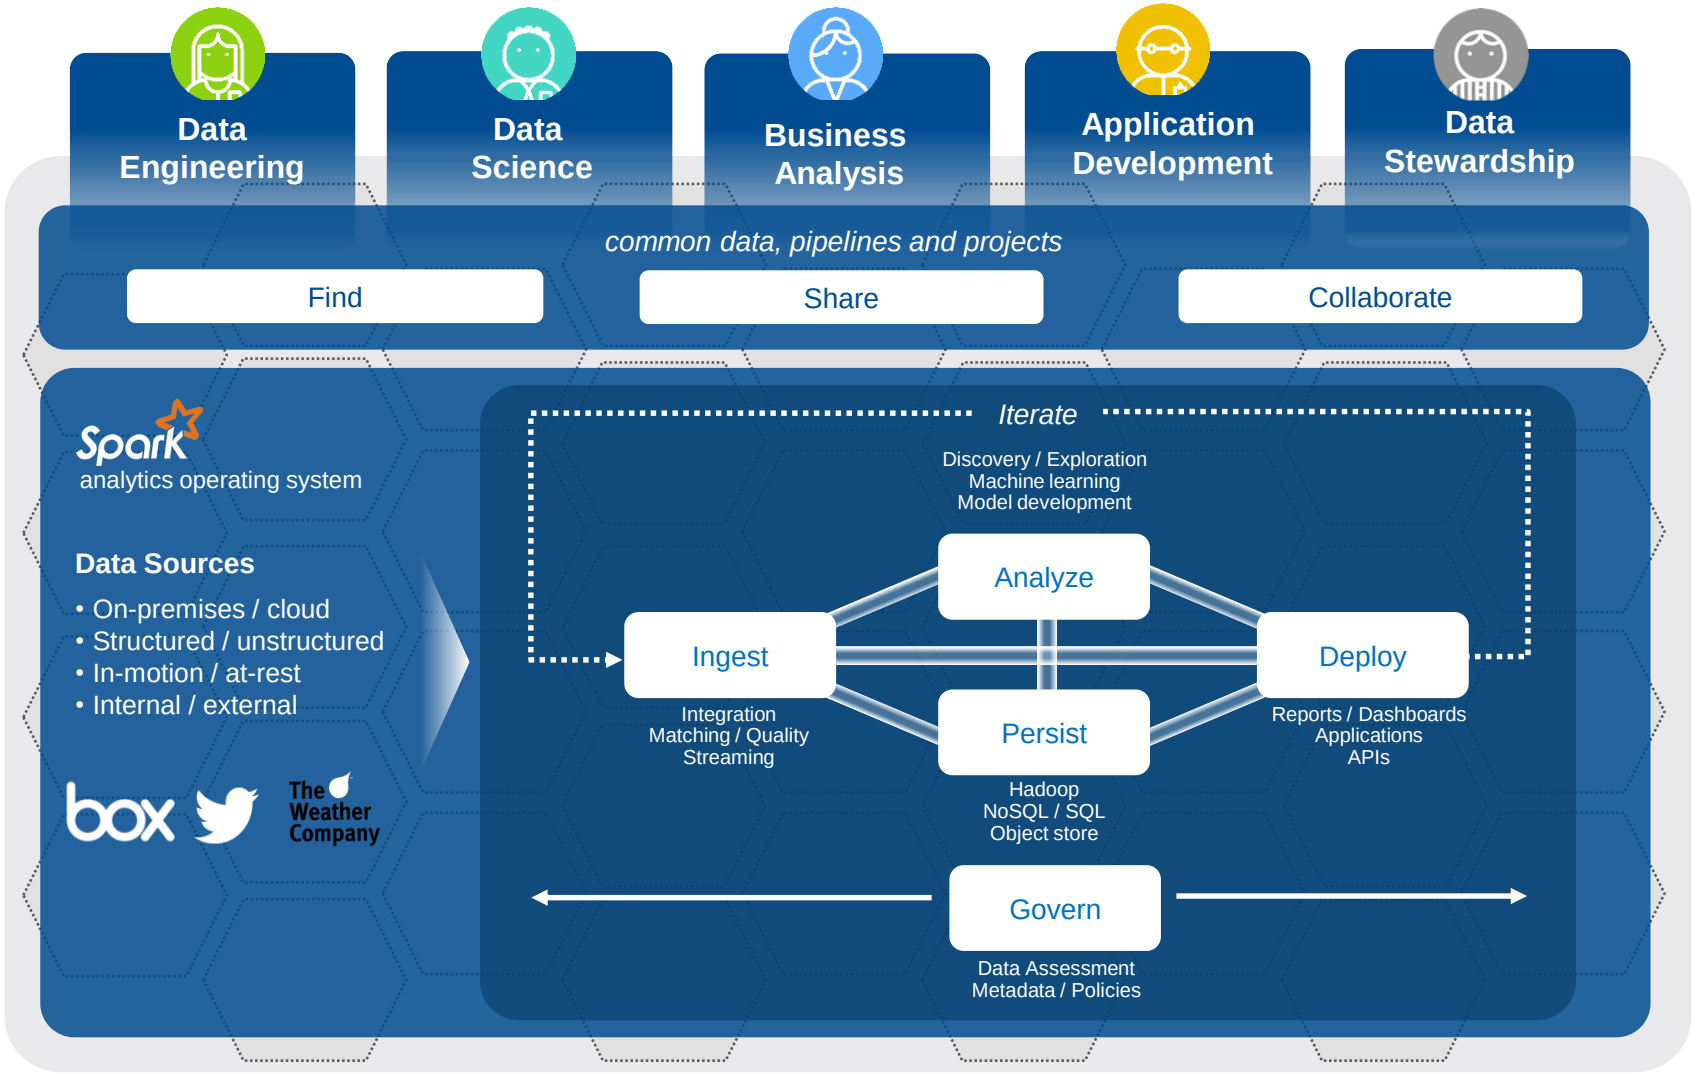

Data Science Experience

Data Scientist

Brings together everything a Data Scientist needs to be successful

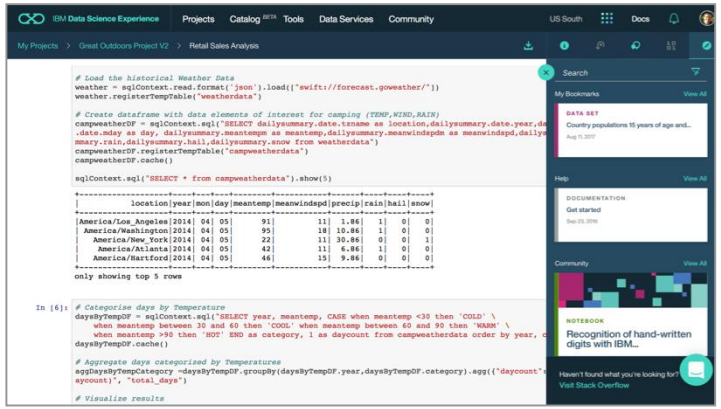

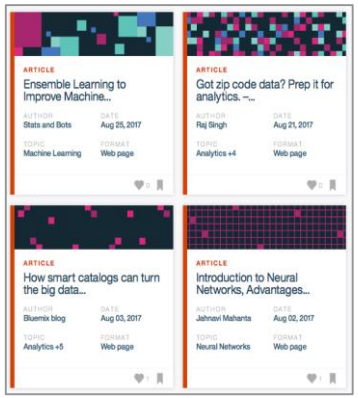

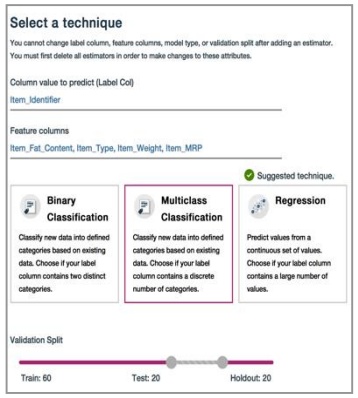

Learn

Built-in learning to get started or go the distance with advanced tutorials

Create

The best of open source and IBM value-add to create state-of-the-art data products

Collaborate

Data and Analytic assets are contained within projects which can be shared with other users.

IBM Cloud PaaS

Rich Platform and Service APIs for your developers

Application Developer

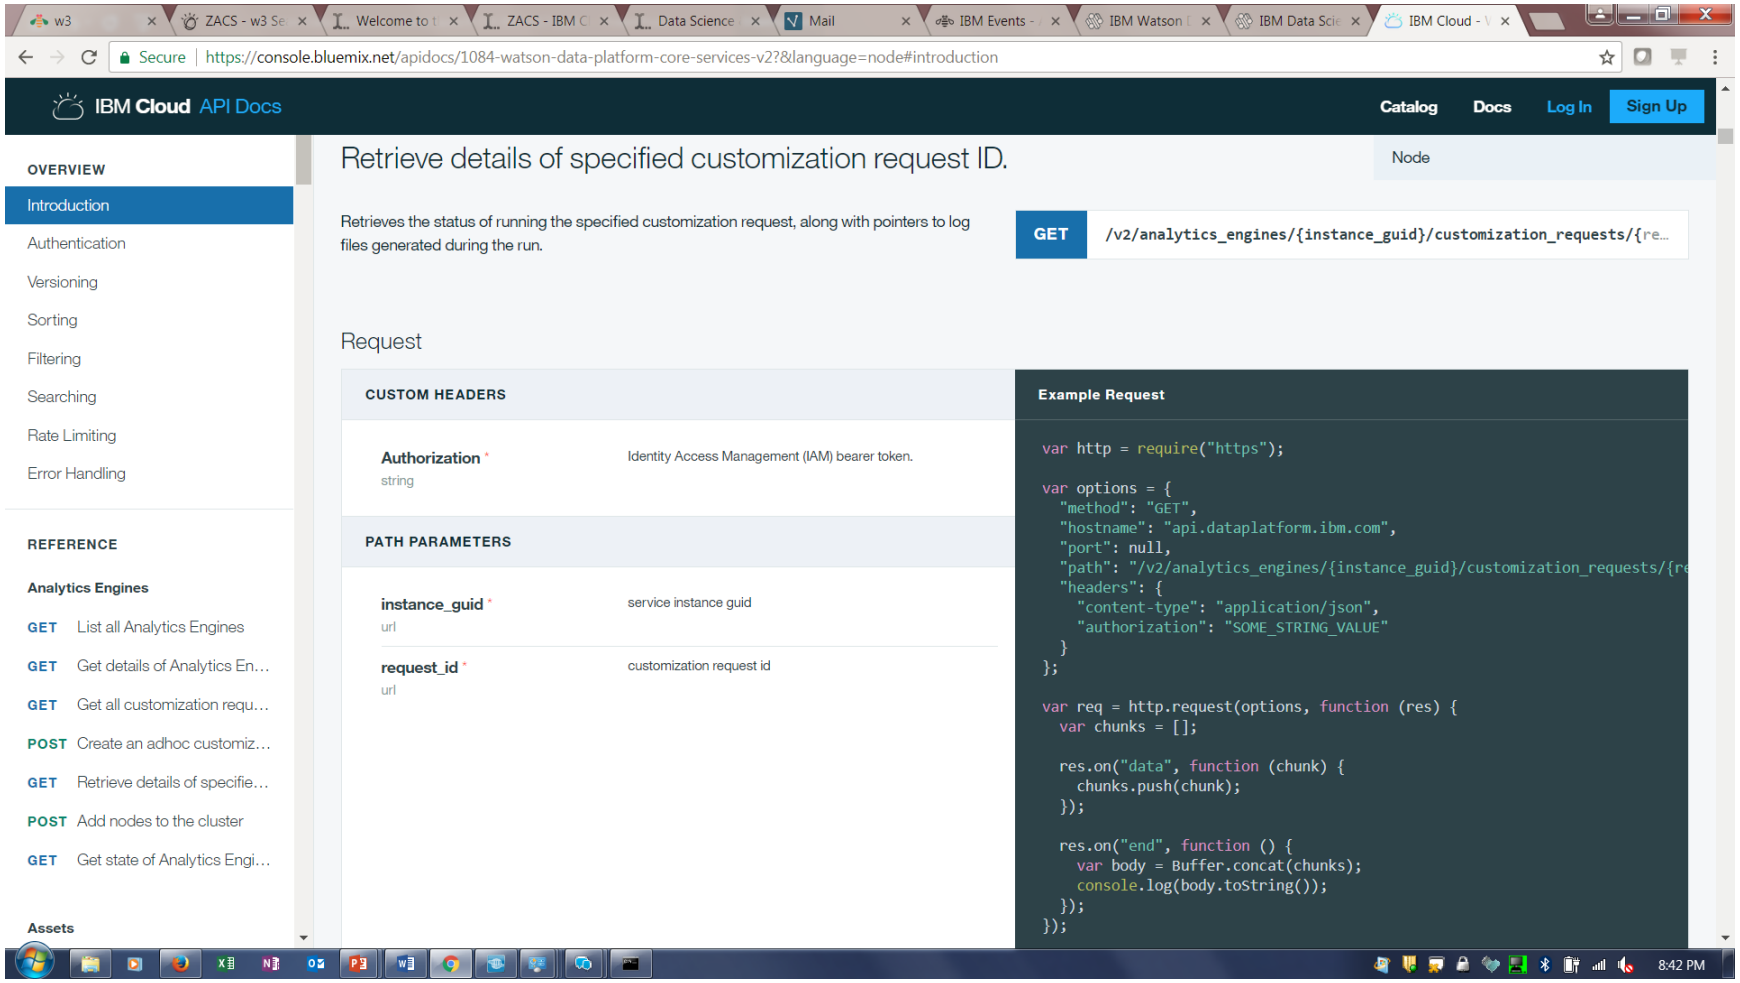

Intelligent data fabric provides consistent platform experience

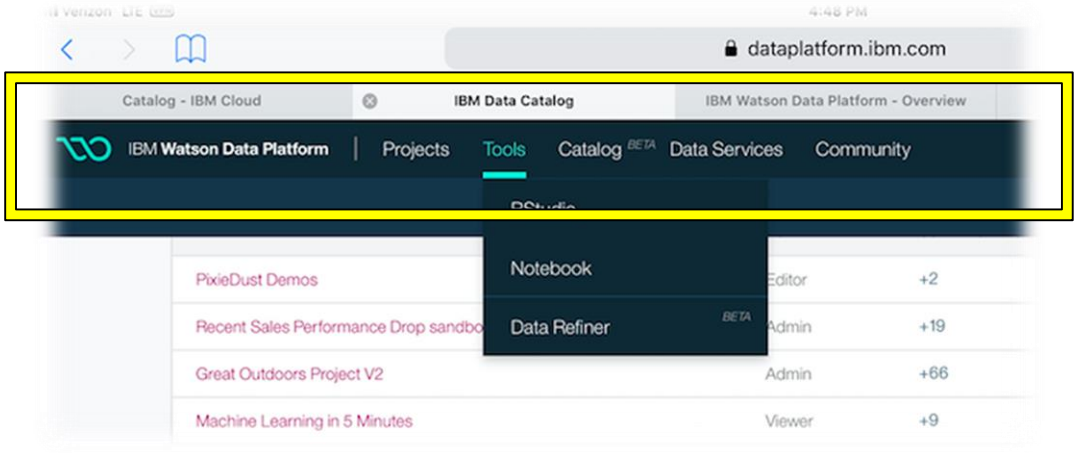

This fabric remains consistent throughout the Watson Data Platform experience – regardless if you are ingesting data, shaping data, building algorithms, deploying models and more…

How does WDP help fulfill the promise of your data?

Data

Puts every important data source at the fingertips of the teams that need it wherever resides

Governance

Enforces your policies without getting in the way of delivering insights

Skills

Makes the most of the data professionals you have and helps them grow and learn from each other as a team

Infrastructure

Delivers the foundation for your first data project through to the complete transformation of your business

Data Science Experience

Core Attributes of the Data Science Experience

IBM Data Science Experience

Community	Open Source	IBM Added Value
• Find tutorials and datasets	• Code in Scala/Python/R/SQL	• IBM Machine Learning
• Read articles and papers	• Jupyter Notebooks	• SPSS Modeler Canvas
• Connect with Data Scientists	• RStudio IDE and Shiny	• Prescriptive Analytics - DOcplexcloud
• Share comments	• Apache Spark	• Projects and Version Control
• Copy and share notebooks	• Your favorite libraries	• Managed Spark Service

Powered by IBM **Watson Data Platform**

DSX Architecture

DSX architecture

Last updated: June 27, 2017

DSX provides you with the environment and tools to solve your business problems by collaboratively analyzing data. This illustration shows how the architecture of DSX is centered around the project. A project is how you organize your resources for solving a business problem.

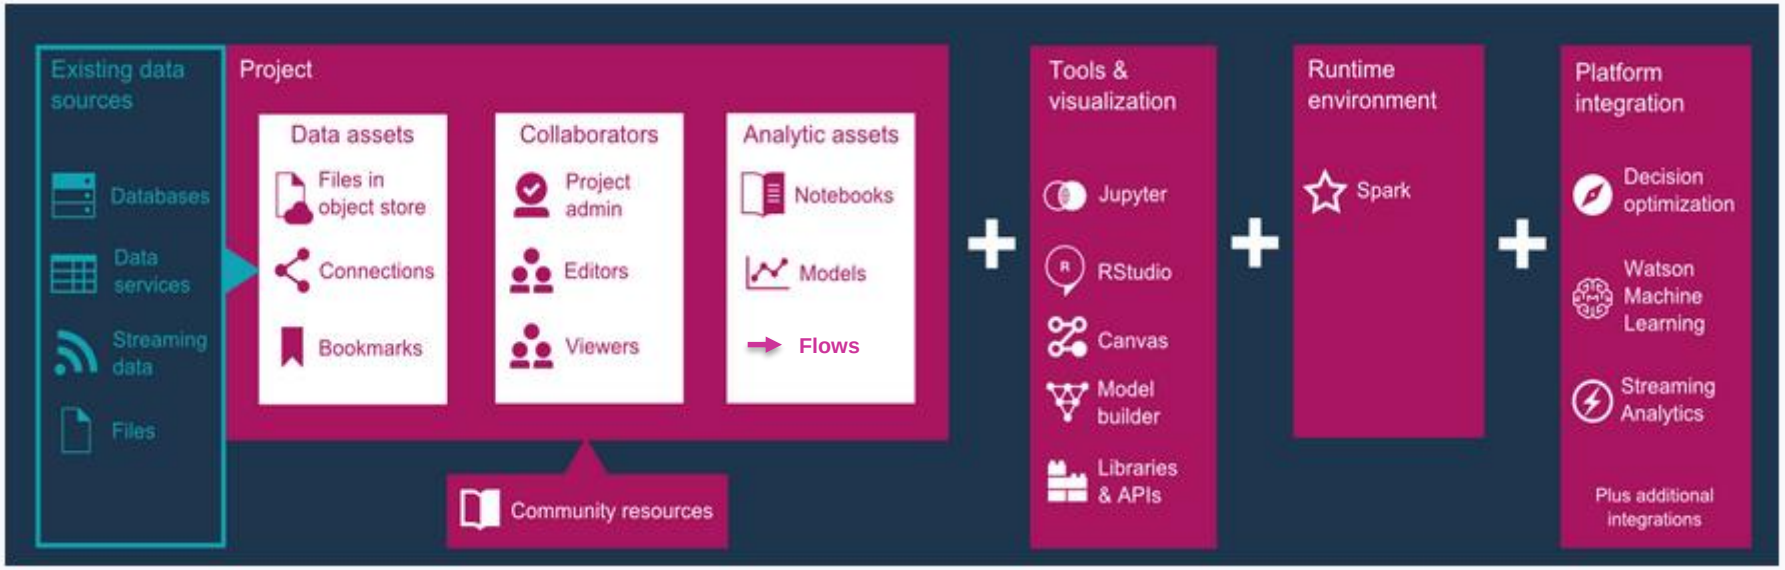

Community Cards provide in-context learning

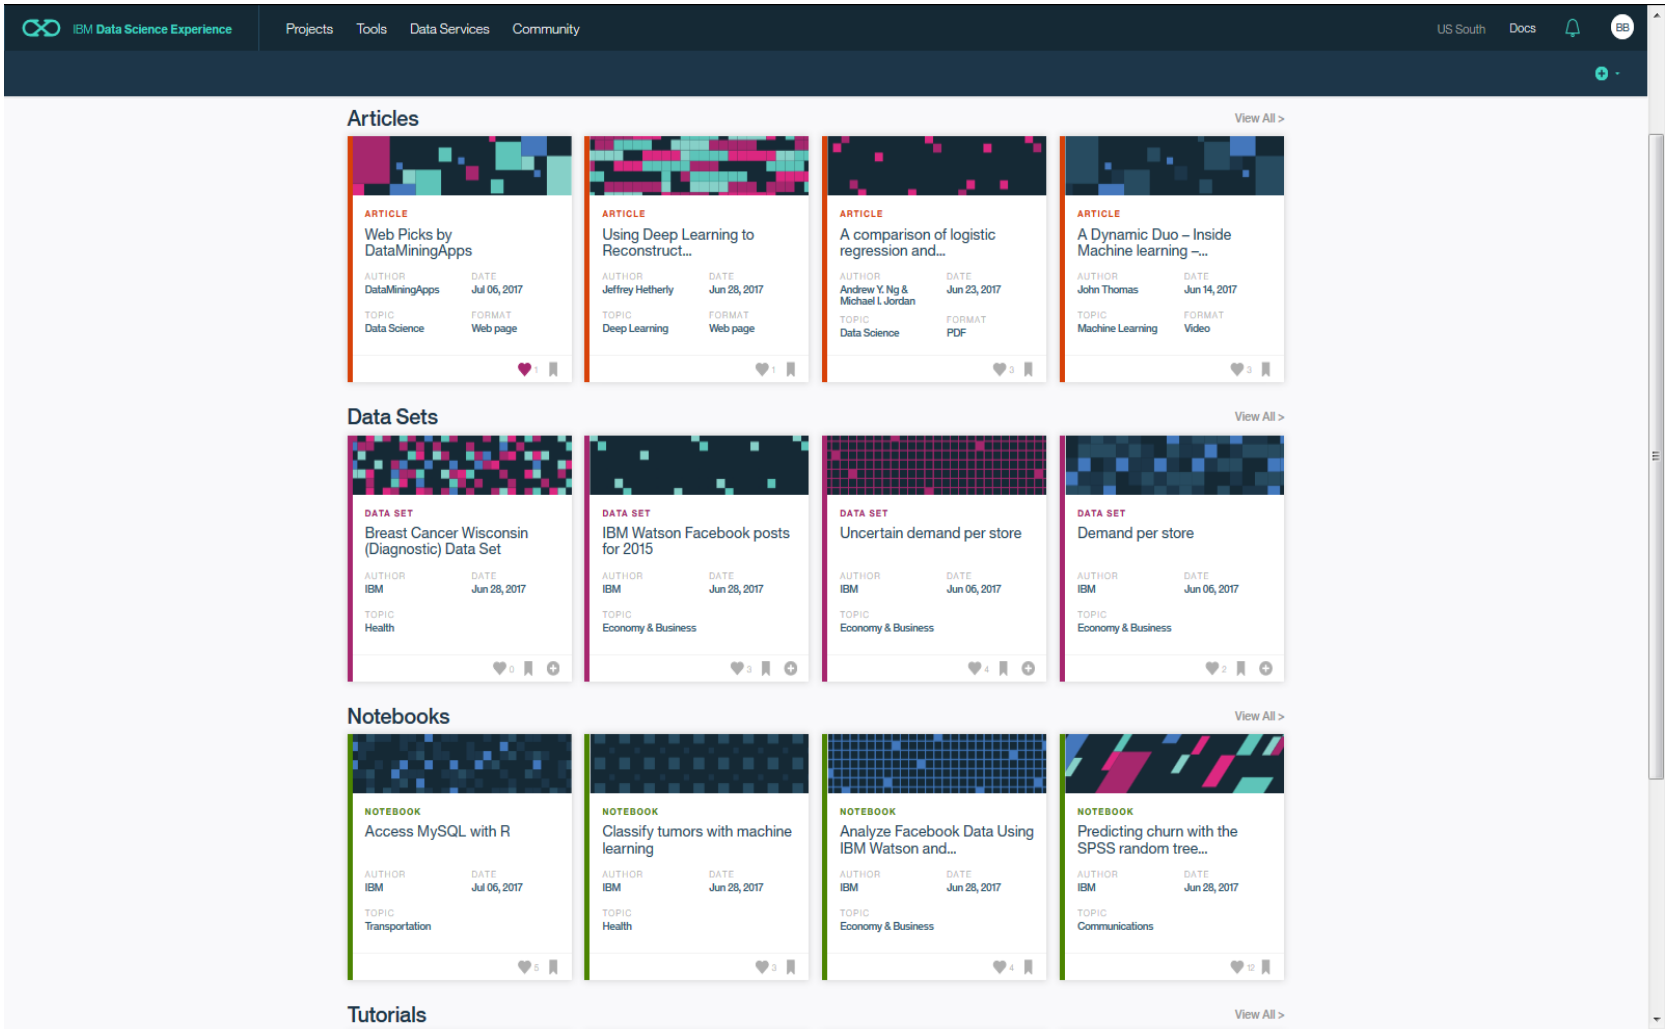

Collaborate Using Projects

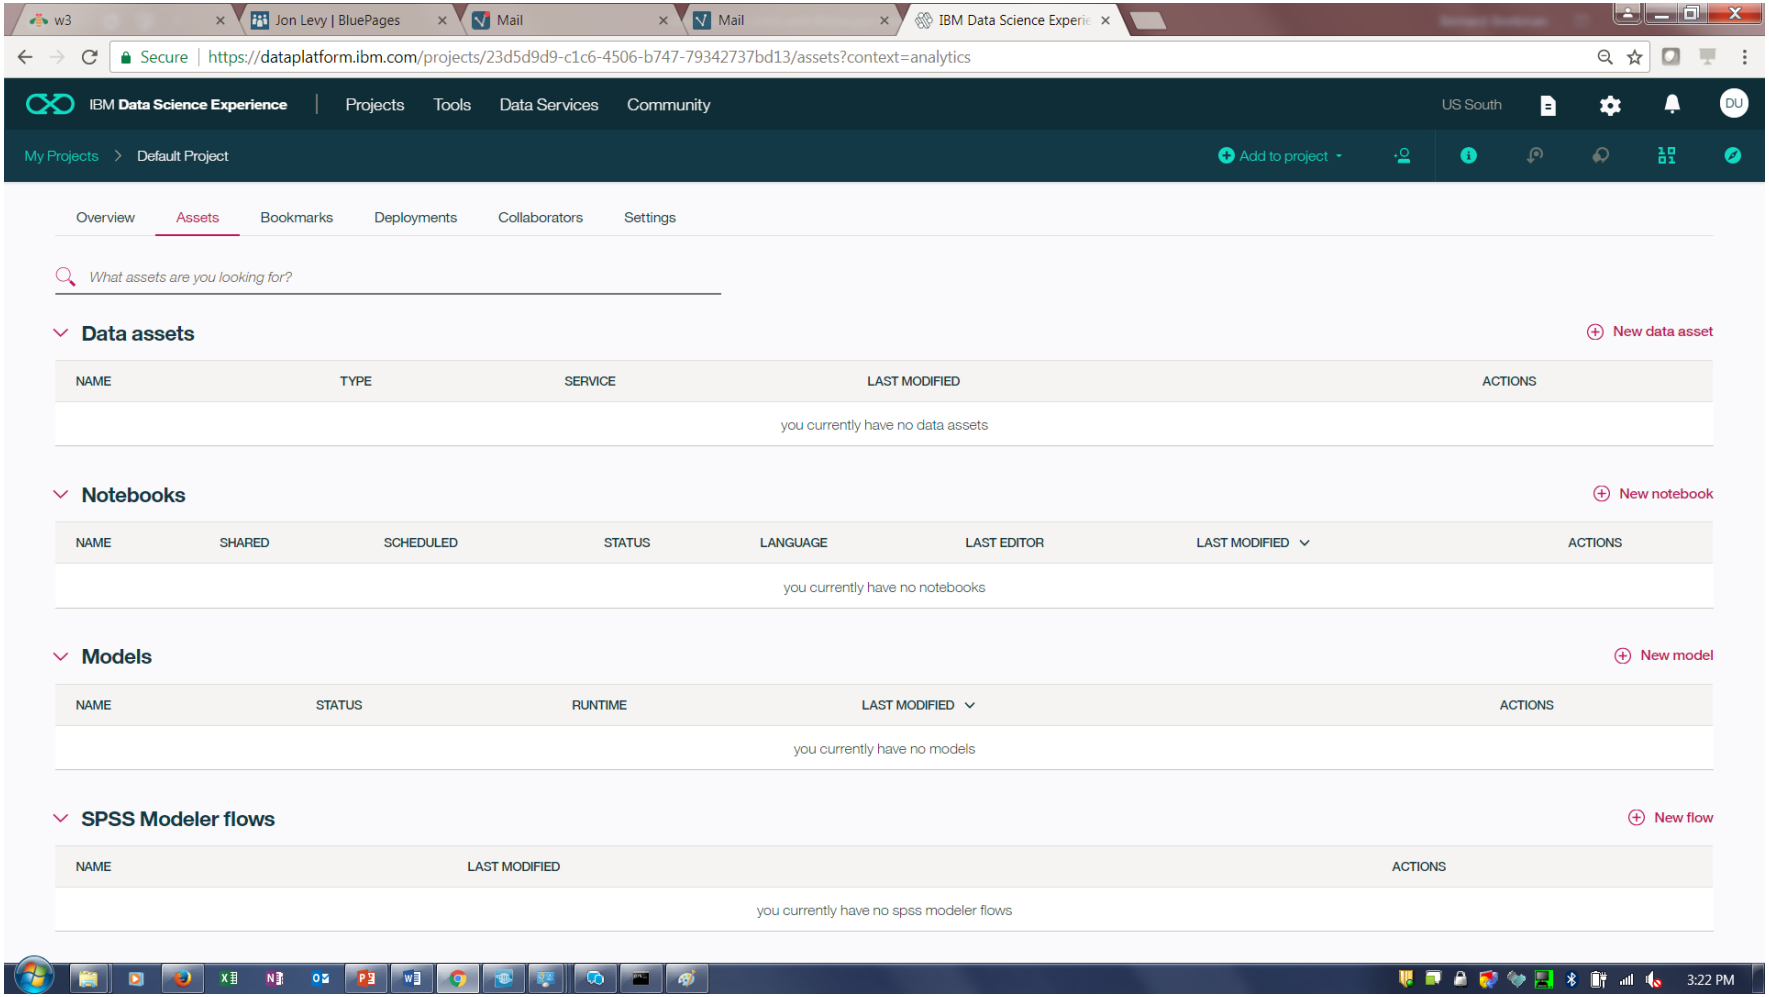

Add Collaborators to a Project

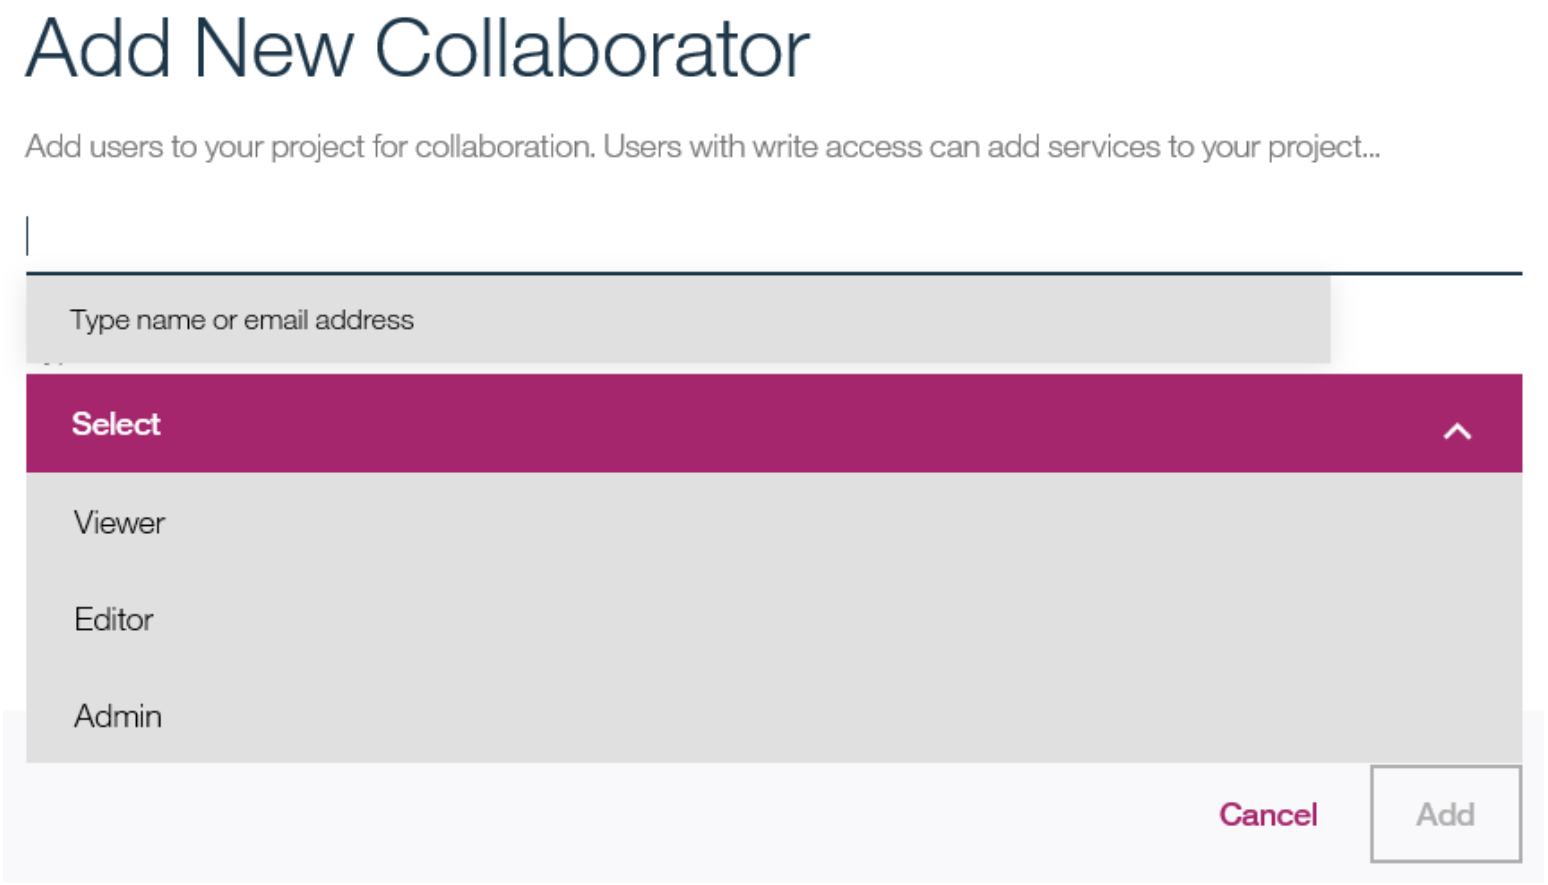

GitHub Integration

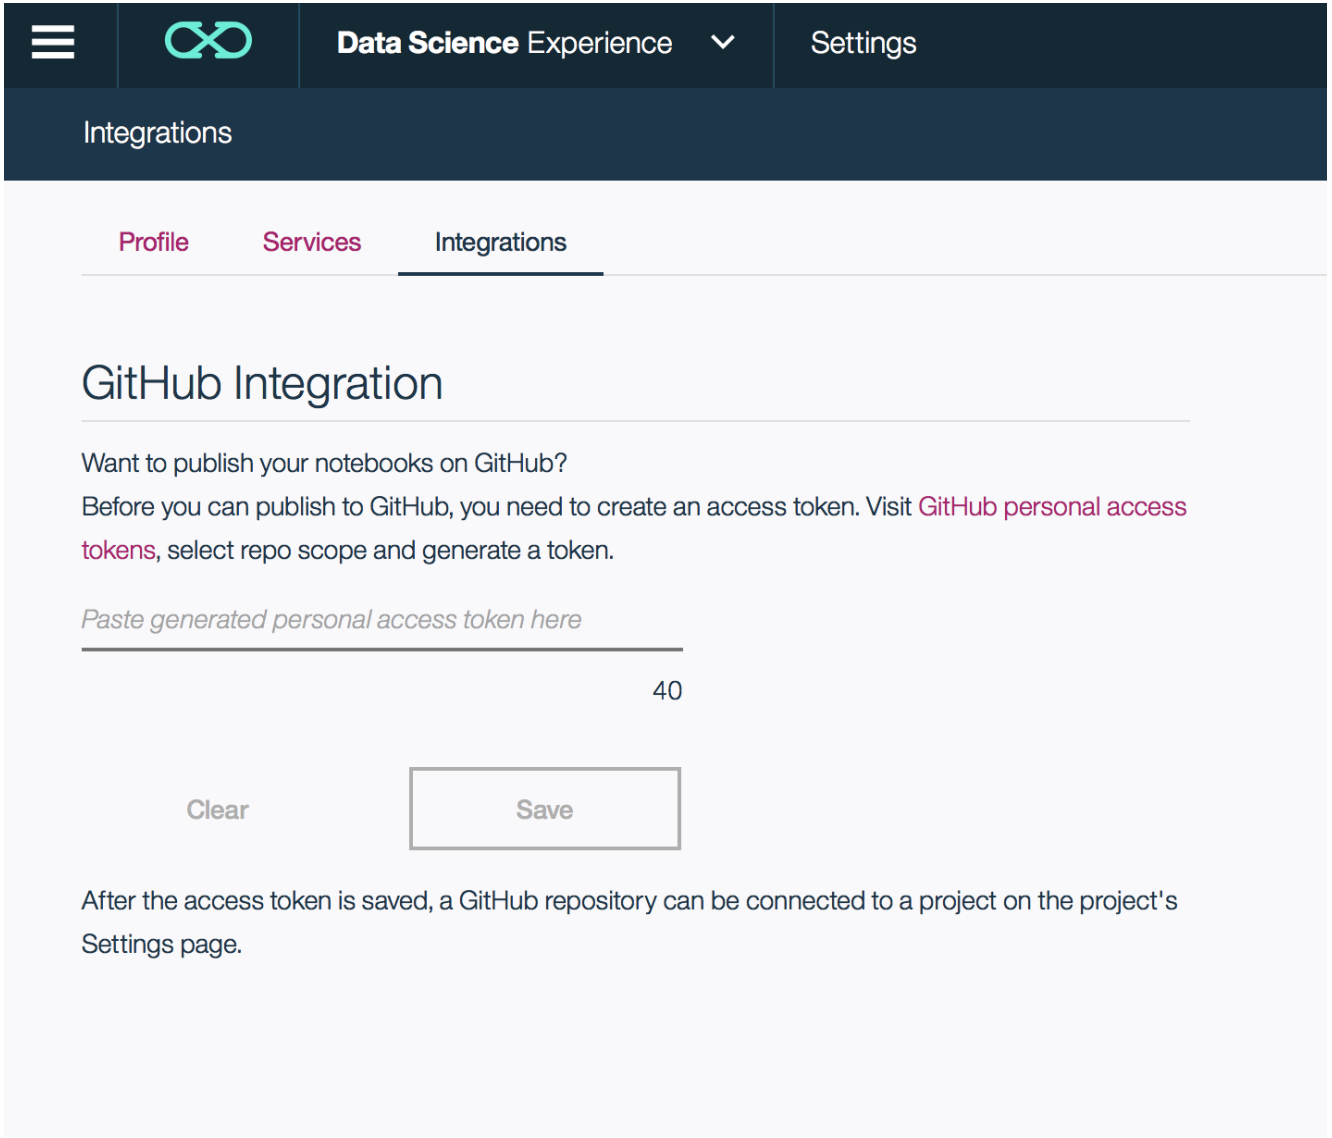

Live chat on Intercom for support from the IBM team and to provide your feedback on how we can improve

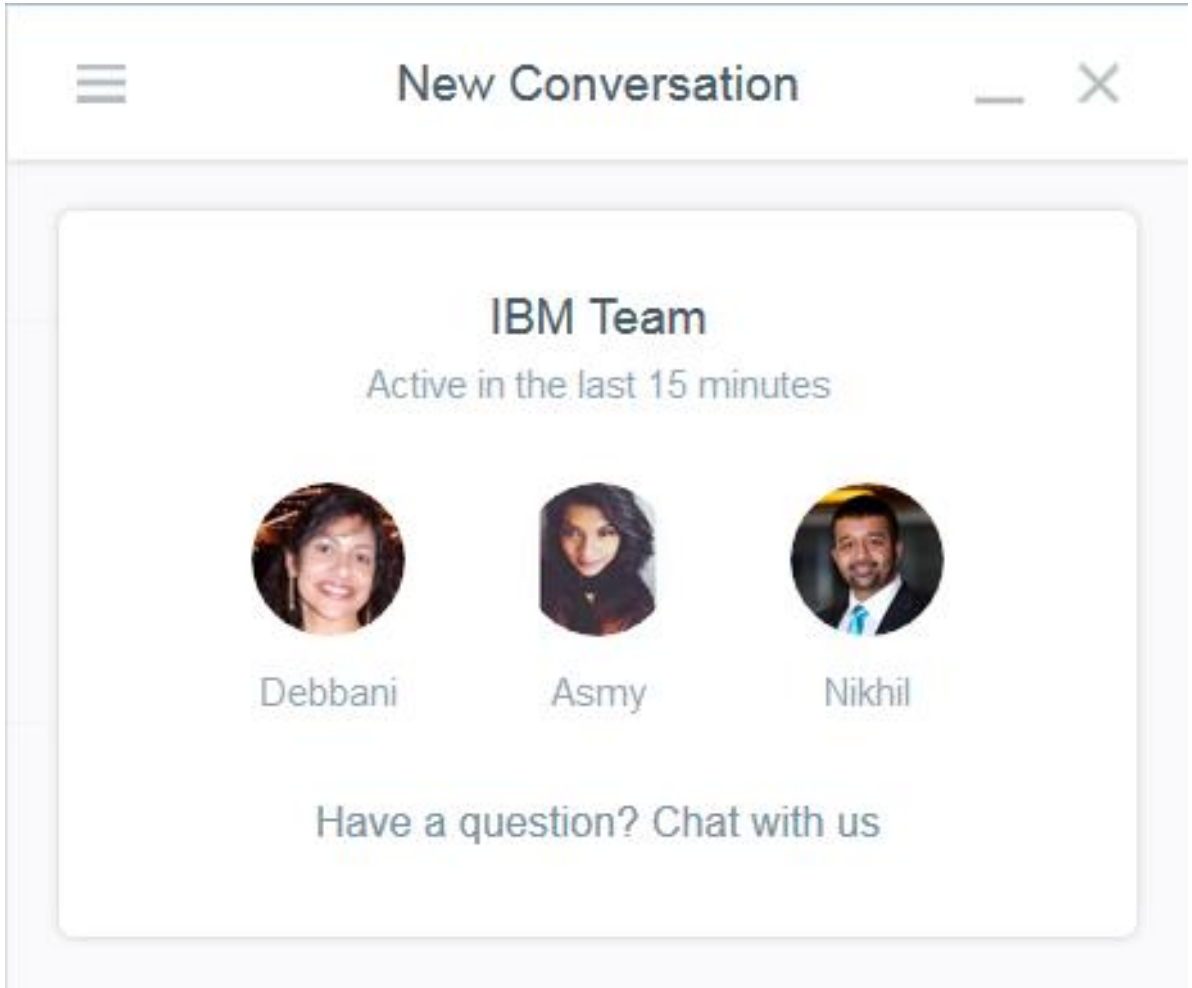

What is a “Notebook”?

Pen and Paper

- **Pen and paper has long provided the rich experience that scientists need to document progress through notes and drawings:**
  - Expressive
  - Cumulative
  - Collaborative

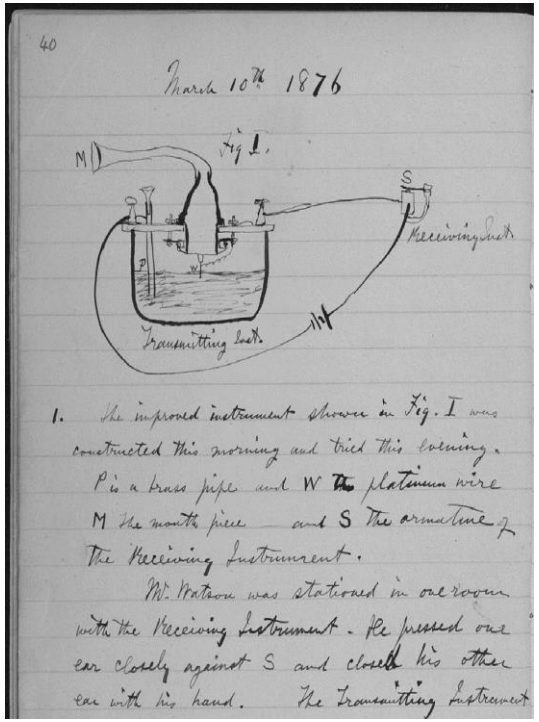

Notebooks

- **Notebooks are the digital equivalent of the “pen and paper” lab notebook, enabling data scientists to document reproducible analysis:**
  - Markdown and visualization
  - Iterative exploration
  - Easy to share

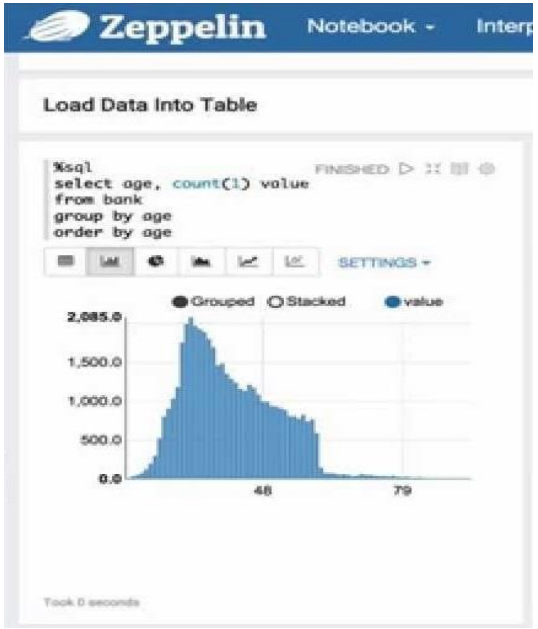

Integrated Jupyter Notebooks for interactive and collaborative development - seamless execution on Spark

Notebooks > Spark Summit West - BlocPower Demo

File Edit View Insert Cell Kernel Help Python 2

Format Markdown Cell Toolbar None

Explore the correlations

In [9]:
```
import matplotlib.pyplot as plt
%matplotlib inline
plt.style.use('ggplot')
ff = pd.tools.plotting.scatter_matrix(dfScaled, diagonal='hist',figsize=(10,10))
```

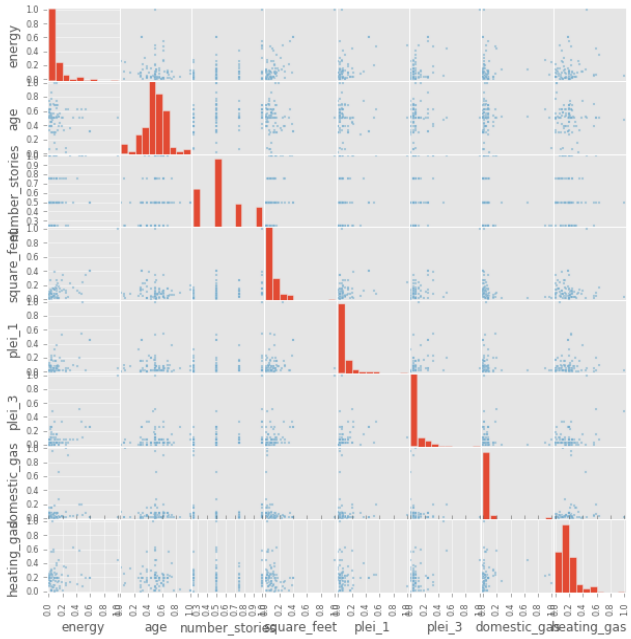

From a Notebook in DSX you can use IBM's managed Spark Service to blend multiple data types, sources, and workloads

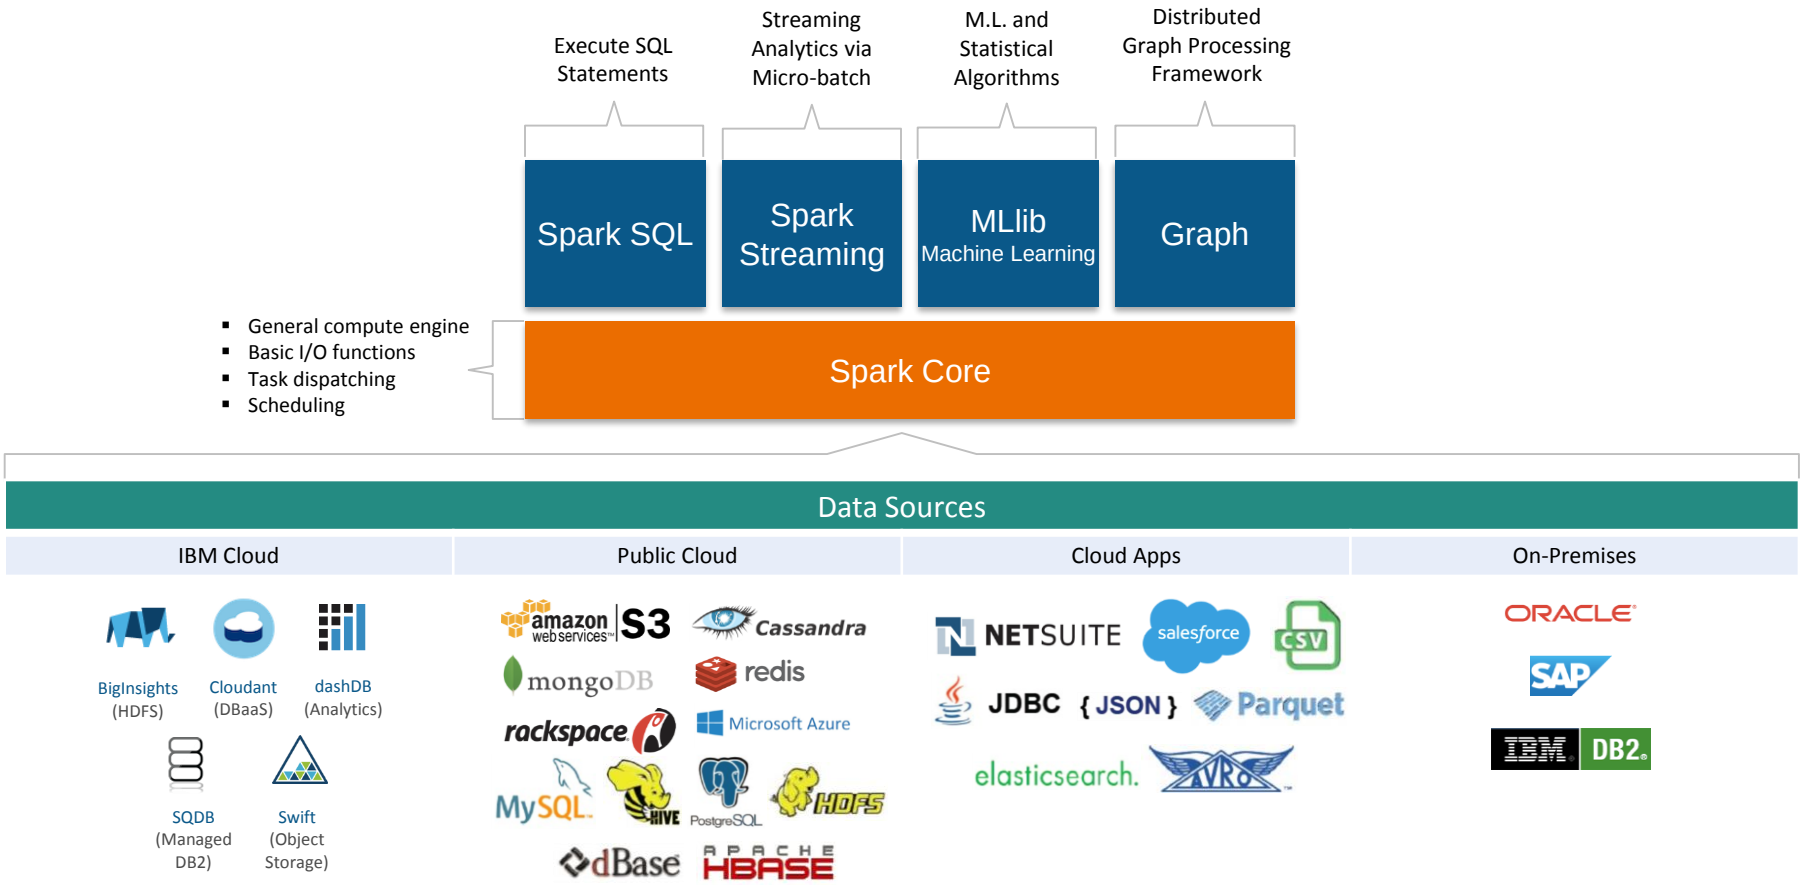

Benefits of Spark for Data Science

- General compute engine
- Basic I/O functions
- Task dispatching
- Scheduling

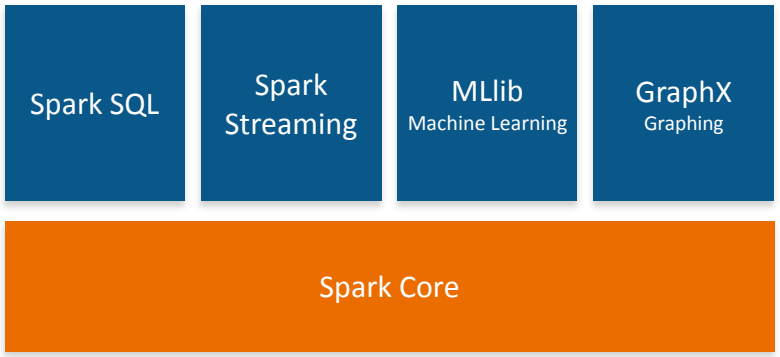

- Allows Data Scientists to code at scale
  - In-Memory processing that scales in a distributed architecture
- Supports multiple programing interfaces (Scala, Python, Java and R)
- Provides unified APIs (SQL, Streaming, Machine Learning, etc.)

The Spark service uses Bluemix Object Storage as its preferred data store for building performant applications

- Object storage provides inexpensive, scalable and self-healing retention of massive amounts of unstructured data
- Every object exists at the same level in a flat address space
- Bluemix Object Storage has a drag-and-drop upload and Swift API for programmatic access

Object Storage

IBM

Supported Data Sources via on- premises and cloud Connectors

IBM services in IBM Cloud

- IBM Informix
- IBM Db2 for i
- Object Storage OpenStack Swift for IBM Cloud
- Object Storage OpenStack Swift (IaaS)
- PostgreSQL on Compose
- IBM Cloudant
- IBM Db2
- IBM PureData for Analytics
- MySQL on Compose
- Cloud Object Storage (IaaS)
- IBM BigInsights HDFS
- IBM Db2 for z/OS
- Cloud Object Storage
- IBM Db2 on Cloud
- IBM Db2 Hosted
- IBM Db2 Warehouse on Cloud

Third-party services

- Cloudera Impala
- Microsoft SQL Server
- Amazon S3
- Pivotal Greenplum
- Salesforce.com
- Sybase IQ
- MySQL
- Microsoft Azure SQL Database
- Apache Hive
- Sybase
- Hortonworks HDFS
- Amazon Redshift
- Oracle
- PostgreSQL

DSX has RStudio built into the experience thanks to our strategic partnership

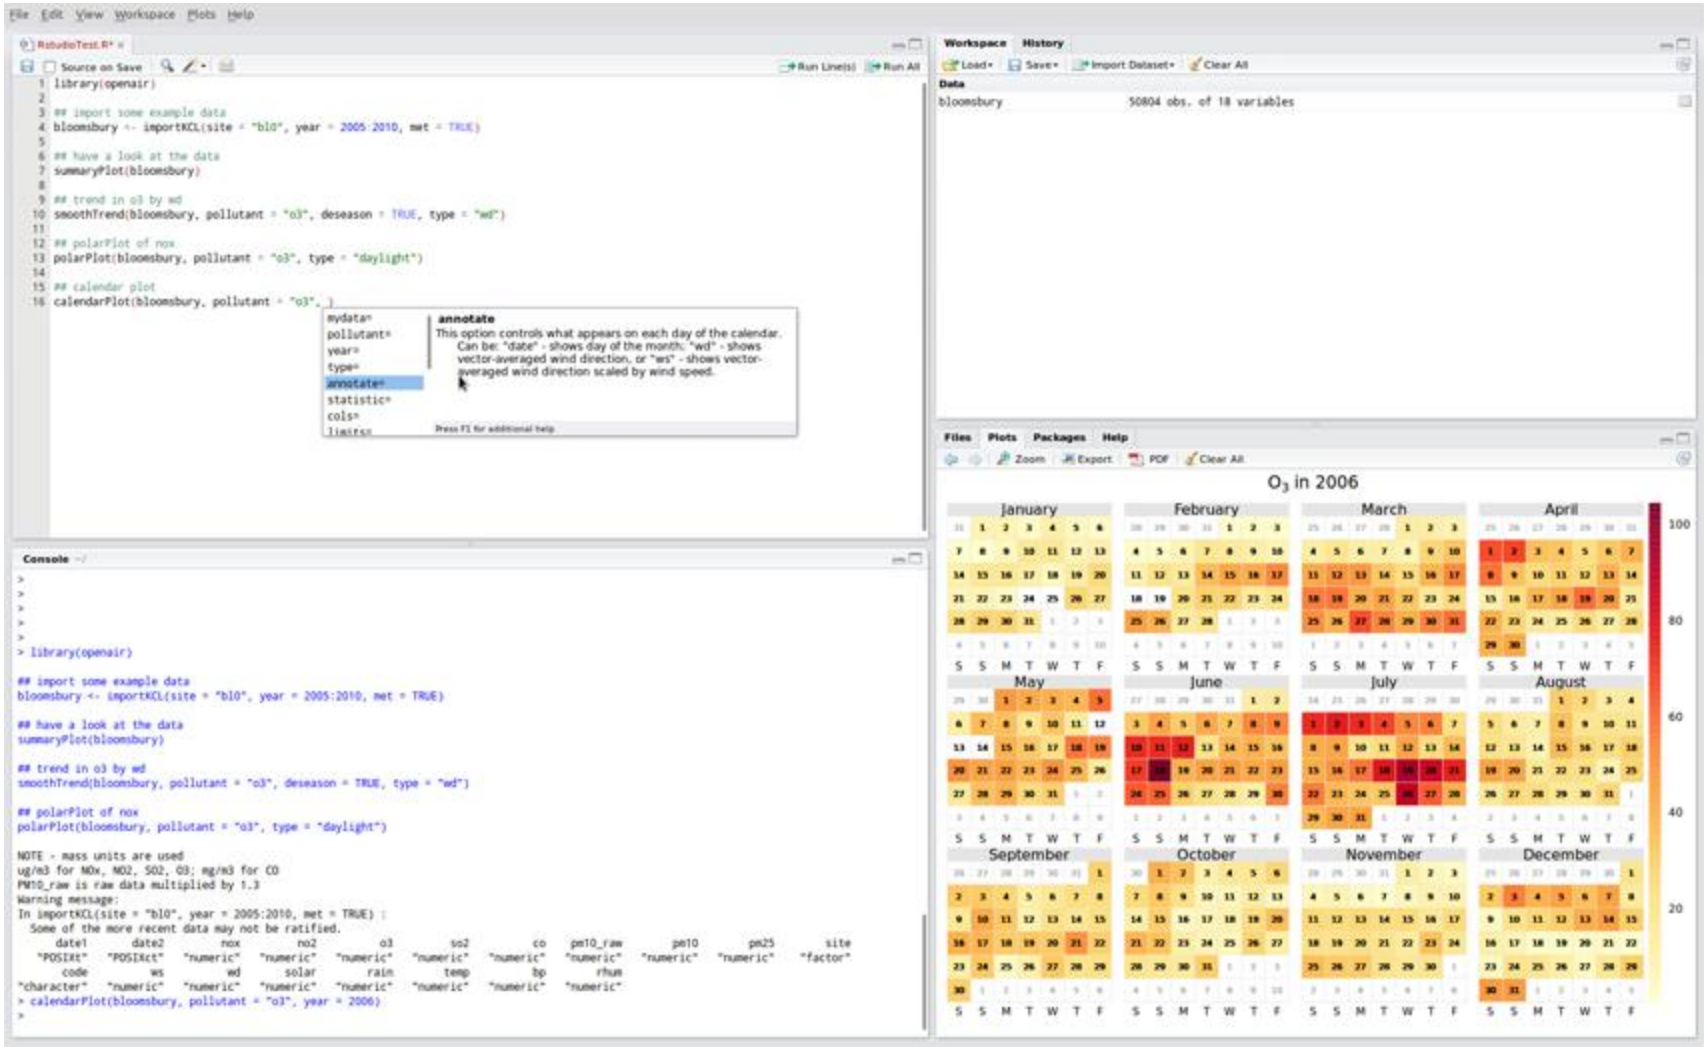

With RStudio you can create Shiny web applications to make your analysis accessible to the business

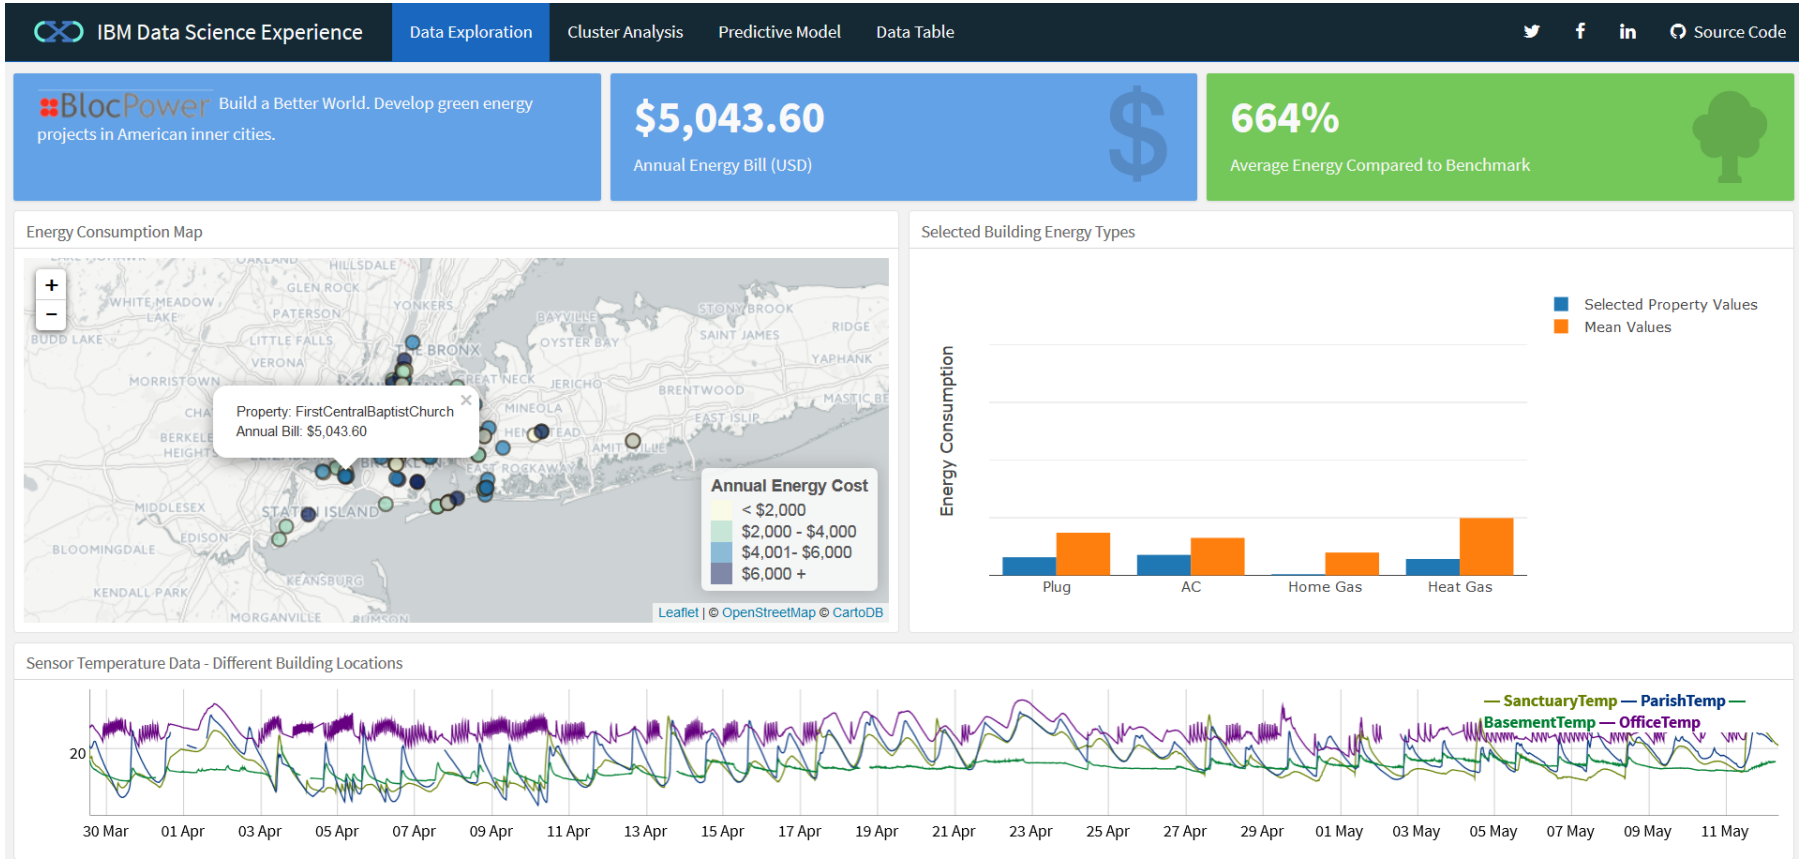

Operationalize insights with Machine Learning

Watson Machine Learning

Data Access:

- Easily connect to Behind-the-Firewall and Public Cloud Data
- Catalogued and Governed Controls through Watson Data Platform

Creating Models:

- Single UI and API for creating ML Models on various Runtimes
- Auto-Modeling and Hyperparameter Optimization

Web Service:

- Real-time, Streaming, and Batch Deployment
- Continuous Monitoring and Feedback Loop

Intelligent Apps:

- Integrate ML models with apps, websites, etc.
- Continuously Improve and Adapt with Self-Learning

Use SPSS Modeler to Visually Create ML Flows

- This DSX Canvas will have compatibility with legacy SPSS Modeler streams
- Multiple execution runtimes: SPSS Modeler, SparkML
- Planned support for R/Python/SQL code

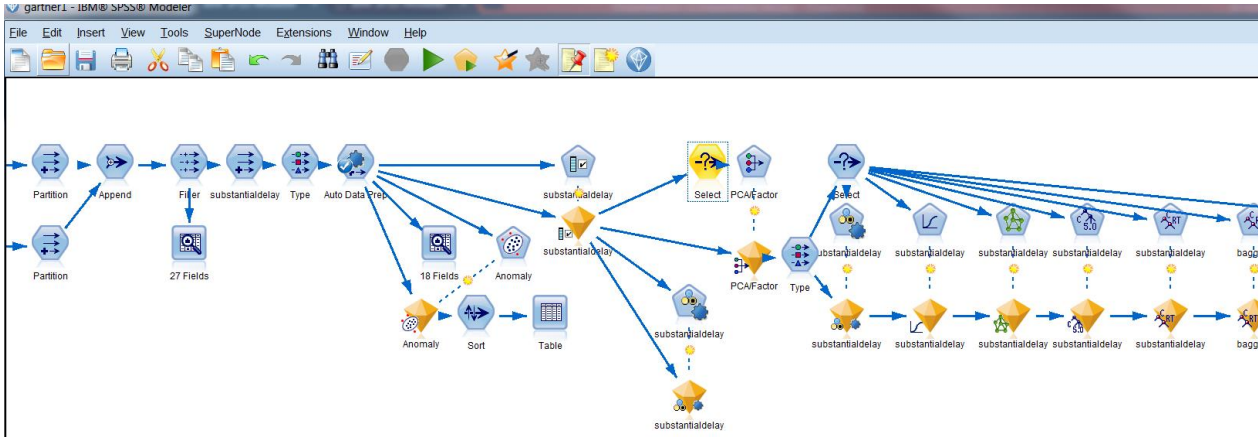

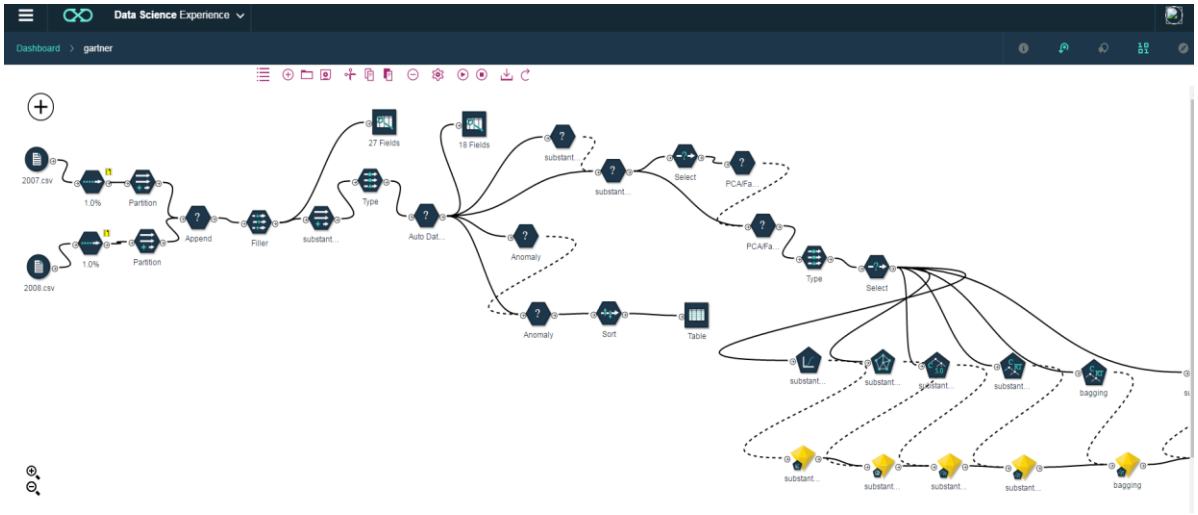

- Pipeline deployment from DSX Canvas (left) via SPSS Modeler

DSX Local

- **Very similar to the public cloud version of DSX**
- **Runs on hardware that is provided by the customer**
  - The DSX Local software and hardware are managed by the customer
- **DSX Local comes with all the software it needs to run, although it can integrate with existing customer systems such as**
  - Databases and HDFS storage
  - LDAP servers for authentication

Labs

Lab Overview

Use IBM's Data Science Experience (DSX) and IBM cloud services to create a working cloud-based application from start to finish. Participants will be led through a series of three labs. The three labs build upon one another so it is important that they are completed in order.

- Lab 1 - The first lab will begin with loading raw delimited data into DB2 Warehouse for Cloud and interacting with that data from a Jupyter notebook in DSX with python.
- Lab 2 - The second lab will guide participants in creating an R notebook and Shiny UI in DSX using RStudio.
- Lab 3 - The third lab will show how to use the Watson Machine Learning capability to create a machine learning model based on the supply chain data set. The machine learning model, deployed in the IBM Cloud, will be used to predict the severity level of each discrepancy based on action request characteristics.

Lab 1

This lab will begin with loading raw delimited data into DB2 Warehouse for Cloud and interacting with that data from a Jupyter notebook in DSX with Python.

Objectives:

- Upon completing the lab, you will know how to:
  - Create a Jupyter IPython notebook from a URL
  - Establish a connection to DB2 Warehouse on Cloud
  - Use a dataframe to read and manipulate tables
  - Use Spark to explore and analyze the dataset
  - Write the modified dataset back to DB2 Warehouse on

Lab 2

In this lab, you will learn some of the fundamentals of using RStudio and Shiny in DSX to work and interact with data in DB2 Warehouse and then create a fully operational "reactive" web application that you can enhance further.

Objectives:

- Upon completing the lab, you will:
  - Create an RStudio project from a Git repository
  - Establish a connection to DB2 Warehouse
  - Query, explore and visualize data in an R notebook
  - Use ggplot2 to create bar plots of several columns in an R dataframe
  - Close the database connection
  - Leverage shiny to create and run a web application
  - Interact with the shiny web application by running it externally

What is Machine Learning?

"Computers that learn without being explicitly programmed"
"Using algorithms to understand patterns in data"

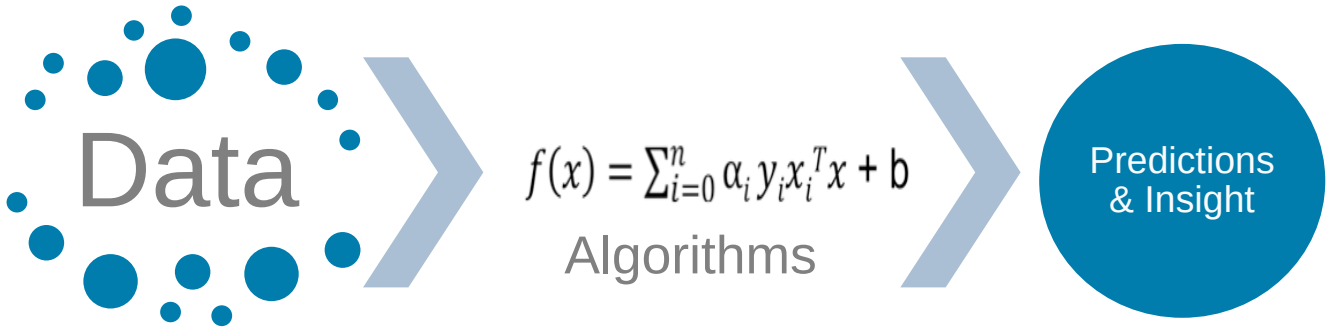

Categories of Machine Learning

- **Supervised learning**
  - The program is “trained” on a pre-defined set of “training examples”, which then facilitate its ability to reach an accurate conclusion when given new data
  - The algorithm is presented with example inputs and their outcomes (labels)
  - The goal is to learn a general rule that maps inputs to outputs

- **Unsupervised learning**
  - No labels are given to the learning algorithm, leaving it on its own to find structure (patterns and relationships) in its input

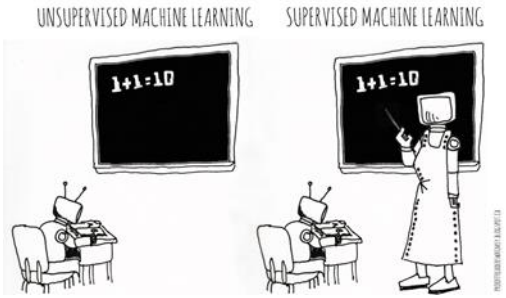

Categories of Machine Learning

Technique	Usage	Algorithms
Classification (or prediction)	• Used to predict group membership (e.g., will this employee leave?) or a number (e.g., how many widgets will I sell?)	• Decision Trees • Logistic Regression • Random Forests • **Naïve Bayes** • Linear Regression • Lasso Regression etc
Segmentation	• Used to classify data points into groups that are internally homogenous and externally heterogeneous. • Identify cases that are unusual	• K-means • Gaussian Mixture • Latent Dirichlet allocation etc
Association	• Used to find events that occur together or in a sequence (e.g., market basket)	• FP Growth

Training, testing, & validation sets

- **During the model development process, supervised learning techniques employ training and testing sets and sometimes a validation set.**
  - Historical data with known outcome
  - Data is randomly split into training, testing, and/or validation sets (mutually exclusive records)
- **Why?**
  - Training set
    - Build the model
    - Tune the parameters
  - Testing set
    - Assess model quality during training/tuning process
    - Avoid overfitting the model to the training set
  - Validation set
    - Estimate accuracy or error rate of model after tuning
    - Used to compare multiple models

Spark ML

- **Spark ML is Spark's machine learning (ML) library**
- **Goal is to make machine learning scalable and easy**
  - No need to understand the detailed math!
- **Divides into two packages:**
  - spark.mllib contains the original API built on top of RDDs
  - spark.ml provides higher-level API built on top of DataFrames for constructing ML pipelines
  - A pipeline is a series of stages where each stage either transforms, or runs through a machine learning algorithm.
- **Using spark.ml is recommended because with DataFrames the API is more versatile and flexible**
  - spark.mllib will continue to be supported

Spark ML Pipeline Terminology

Spark ML standardizes APIs for machine learning algorithms to make it easier to combine multiple algorithms into a single pipeline, or workflow

- **DataFrame**: Spark ML uses DataFrame from Spark SQL as an ML dataset, which can hold a variety of data types
- **Transformer**: A Transformer is an algorithm which can transform one DataFrame into another DataFrame
- **Estimator**: An Estimator is an algorithm which can be fit on a DataFrame to produce a Transformer
- **Pipeline**: A Pipeline chains multiple Transformers and Estimators together in a sequence to specify an ML workflow
- **Parameter**: All Transformers and Estimators share a common API for specifying parameters

Lab 3 - Watson Machine Learning

In this lab, you will use IBM's Watson Machine Learning GUI to train, evaluate, and deploy a Watson Machine Learning model based on the modified supply chain dataset.

Objectives:

- Upon completing the lab, you will:
  - Become familiar with the Watson Machine Learning GUI.
  - Train/Evaluate a machine learning model
  - Deploy a machine learning model.
  - Use the deployed machine learning model to make predictions.